




Designed boundary fixtures with author-chosen thresholds. The values are not measurements of any deployed model, and the $n = 18$ fixture sits where percentile-bootstrap bias is not negligible. The altered-profile fixture has scores identical to the supported fixture: its denial is driven by the profile digest alone. The four plotted fixtures cover 90 of the 104 transactions; the remaining 14 use two profile variants this laboratory owns in the same fixture schema, scored identically to the supported fixture and passed by the unmodified SAFE evaluator, which differ only in profile identity and are therefore not SAFE boundary cases.